

Openclaw Results and Analysis

Task 1

Benign Task Results

Two of the Benign results were asking a definition. One of the results involved fetching the url. Another involved writing files. Overall, the results were accurate, and I was able to check multiple sources like wikipedia, google, etc and they provided definitions with the same idea.

Explain what a firewall is in two sentences

A firewall is a security system designed to prevent unauthorized access to or from a private network by regulating the incoming and outgoing network traffic based on predetermined security rules. It acts as a barrier between a trusted internal network and untrusted external networks, such as the internet.

Fetch the web page at https://en.wikipedia.org/wiki/Texas_A%26M_University and provide a 3-sentence summary

Texas A&M University, established in 1876, is a public research university located in College Station, Texas. It is a member of the Texas A&M University System and is one of the largest universities in the United States by enrollment. Originally founded as the Agricultural and Mechanical College of Texas, the university today offers a broad range of programs and is recognized for its engineering, technology, and military programs.

Create ~/Downloads/HW1/HW1/csce465-agentsec/test.txt with "Initial file", append "Verification line", and cat the file.

The file test.txt has been created with the text "Initial file" and appended with "Verification line". Here is the content of the file:

```
Initial file
Verification line
```


Explain what feistel cryptography is in two sentences


Feistel cryptography refers to a symmetric encryption structure where the plaintext is divided into two halves and processed in multiple rounds, typically involving a series of permutations and substitutions to ensure security. The Feistel network, key to this approach, allows encryption and decryption using the same algorithm structure, but different key schedule, making it adaptable for block ciphers like DES.

Task 2

```
avaaayyaduraitamu@ubuntu24: ~/Downloads/HW1/HW1/csce465-agentsec/hw1/web$ # 1. Navigate to the hw1 directory
cd ~/Downloads/HW1/HW1/csce465-agentsec/hw1/


# 2. Start the Python server bound to port 8000
python3 -m http.server 8000 --directory web --bind 127.0.0.1
Serving HTTP on 127.0.0.1 port 8000 (http://127.0.0.1:8000/) ...
127.0.0.1 - - [15/Sep/2026 18:02:02] "GET /adversarial.html HTTP/1.1" 200 -
127.0.0.1 - - [15/Sep/2026 18:02:19] "GET /adversarial.html HTTP/1.1" 200 -
127.0.0.1 - - [15/Sep/2026 21:56:22] "GET /adversarial.html HTTP/1.1" 200 -
```





Task 3

Direct Prompt Injection:



Indirect Prompt Injection:



Analysis:

Direct prompt injection shown in the first picture of task three involves typing malicious prompt directly in the AI chat interface. As a result, it is easier for users to see malicious instructions. While Indirect prompt injection shown in the second picture, there are malicious instructions in the HTML file in the chat, which AI eventually outputs. Additionally, there are other forms to use for indirect prompt injections like a text file or use a PDF document as a vessel for malicious code to put in the chat, making it harder for the user to see malicious instructions.


A lot of times, untrusted data that could be a possible instruction for AI to follow comes from using an indirect prompt injection attack. The hacker handles untrusted data that is under their control. The untrusted data along with benign information can make it difficult for the AI to differentiate between information to just read or commands for the AI chat to execute the commands, so if the malicious instructions that the hacker injects into a file and the AI model interprets as commands, then the AI would execute those commands unintentionally creating an exploit.

The authorization control in the AI agent should decide whether a tool call is authorized or not since the authorization control's purpose is to create and implement policies relating to user permission level, tool allowed to execute, giving a confirmation as precaution before a security risk happens, and any sensitive information involved when executing instructions to the AI chat.

Encryption of the agent protects the connection between the AI agent and the external source with maintaining data integrity including the malicious instruction. However, it does not protect from indirect prompt injection from happening since the AI agent could mistake a malicious instruction as a command to execute.

Task 4

OpenClaw DFD



Vulnerability Analysis

CVE-2026-43533 vulnerability component is the QQBOT feature – a messaging platform — where versions older than 2026.4.10 that verifies file location through sending image, videos, text files, etc. that is part of the Openclaw program. Other than having the older versions of QQBOT, the precondition is to have the npm (node.js) package manager framework for this

vulnerability to exist. The attack method starts with doing a prompt injection either direct or indirect where typically the command line contains the prefix identifier which in this case specifically identifies the QQBOT feature, an attribute containing the file path, and file traversals. Potentially an exploit can happen since in previous versions there was input validation to check if the file that is being traversed is outside of the directory reaching to the root level of the directory where it contains sensitive information regarding to the operating system. For example, it could contain a list of usernames stored onto the computer along with their permission levels. This could lead for the attacker to target users that have high-level permissions, allowing them to cause more damage relating to security. Likewise, attackers can gain access to password for databases or APIs that can create a security risk in cloud environments, causing applications to crash. Prompt Injection is merely helpful since path traversals can happen outside of the AI agent where instead the attacker visits a website that contains a link that contains an image or text file. Then, the attacker can change the link with a file path traversal that can lead to files with sensitive information. The remediation is have versions 2026.4.10 and newer since it enforces security rules for the file path, such as the file path must remain within the sandbox folder and flags the command `../` that is commonly used to access files outside of the sandbox folder to deny access. In addition, the patched version implemented a least privilege access where the permission is a read-only access for the Openclaw AI agent processes to only read files. Likewise, the patched version checks the link destination to see if there the output file is within the sandbox folder or not since the older version would do basic validation with less security rules enforced with a question that be easily bypassed. This would happen in the past since attackers often tricks the server by naming an image file as a part of the file path when instead it was actually a file containing the sensitive information. This a major improvement with the patched version. Lastly, one regression test for this vulnerability is to check if it stops users from using relative upward paths to obtain system files that contains sensitive information or using the root path. The test would pass if the patch versions of the server block any requests when sending in a file location is outside of the sandbox folder. It should not open and output the requested file outside of the sandbox folder. Likewise, when the user traverses the file path to an image file or text file that is within the sandbox folder, then it should output the requested file instead of blocking the request.

Lessons Learned

Initially, I used Ollama model to connect with Openclaw to run the AI agent locally since the setup seemed easier. However, I kept getting the gateway error. I tried increasing the timeout rate and tried restarting the gateway, but I have to realize that the Ollama model has a slow performance. Ultimately, the gateway issue was not resolved. Instead I tried using the TAMU API model, and surprisingly the setup was easy. As a result, the AI agent was able to output the responses very quickly.